

Project Part One

Design Decisions for "hospital" table

The table for "hospital" has the hospital id as its primary key, which is referenced by the "weekly logs" table and the "hospital_quality" table. All hospitals should have a name, so we restrict TEXT NOT NULL for "hospital_name". The "emergency_services" is a BOOLEAN because it indicates whether or not the hospital provided emergency services. Zipcode is a foreign key that references the "locations" table. All other attributes of the hospital -- "address", "geocoded_hospital_address", "type_of_hospital", and "type_of_ownership" -- are text-based and can have NULL values because it is information that may not be available but isn't crucial to have in order for the hospital to be included. Information such as the type of hospital, type of ownership, and emergency services are recorded here, despite being provided in the Hospital Quality data set. This is because we assume that these values are more or less fixed for a given hospital, and will not change very often (if at all).

```
In [ ]: %%sql
CREATE TABLE hospital (
    hospital_pk TEXT PRIMARY KEY,
    hospital_name TEXT NOT NULL,
    address TEXT,
    geocoded_hospital_address TEXT,
    type_of_hospital TEXT,
    type_of_ownership TEXT,
    emergency_services BOOLEAN,
    zipcode INTEGER NOT NULL REFERENCES locations(zipcode)
);
```

Design Decisions for "weekly_logs" table

Each row in the "weekly_logs" table represents one hospital's activity for one week. It has a serial, autoincrementing column "id" as its primary key. It contains the hospital ID as a foreign key which references the 'Hospital' table. This is so that each entry in this table can be linked to/associated with a hospital. It also contains a DATE column which represents the week for which a given entry is from, which we restrict to be NOT NULL, as otherwise a weekly log would be meaningless. This table also contains FLOAT8 columns that represent the number of beds available/in use for adult patients, pediatric care, and the ICU. All of these columns are floating point numeric columns as they are 7 day averages, so it is likely the values will not be integers. They are also constrained to be non-negative, as we can never have less than 0 beds available or less than 0 patients. We also constrain the number of beds in use to be less than or equal to the number of available beds for each category (adult, pediatric, and ICU), as we believe that logically the number of patients should not exceed the number of beds.

In []:

```
%%sql
CREATE TABLE weekly_logs (
    id SERIAL PRIMARY KEY,
    collection_week DATE NOT NULL,
    all_adult_hospital_beds_7_day_avg FLOAT8 CHECK (all_adult_hospital_beds_
    all_pediatric_inpatient_beds_7_day_avg FLOAT8 CHECK (all_pediatric_inpat
    all_adult_hospital_inpatient_bed_occupied_7_day_avg FLOAT8 CHECK (all_ad
    all_pediatric_inpatient_bed_occupied_7_day_avg FLOAT8 CHECK (all_pediatr
    total_icu_beds_7_day_avg FLOAT8 CHECK (total_icu_beds_7_day_avg >=0),
    icu_beds_used_7_day_avg FLOAT8 CHECK (icu_beds_used_7_day_avg >=0),
    inpatient_beds_used_covid_7_day_avg FLOAT8 CHECK (inpatient_beds_used_cc
    staffed_icu_adult_patients_confirmed_covid_7_day_avg FLOAT8 CHECK (staff
    hospital_pk TEXT NOT NULL REFERENCES hospital(hospital_pk),
    CHECK(all_adult_hospital_inpatient_bed_occupied_7_day_avg <= all_adult_h
    CHECK(all_pediatric_inpatient_bed_occupied_7_day_avg <= all_pediatric_in
    CHECK(icu_beds_used_7_day_avg <= total_icu_beds_7_day_avg)
);
```

Design Decisions for "hospital_quality" table

The table for "hospital_quality" stores quality ratings released by CMS for each hospital. We use a surrogate primary key (id) to uniquely identify each record, while hospital_pk serves as a foreign key referencing the "hospital" table to ensure that every rating corresponds to a valid facility. The quality_rating field is numeric. The date_updated column records when the rating was published. We mark both quality_rating and date_updated as NOT NULL because a quality entry must always include a score and a date in order to be meaningful.

```
In [ ]: %%sql
CREATE TABLE hospital_quality (
    id SERIAL PRIMARY KEY,
    quality_rating NUMERIC NOT NULL,
    date_updated DATE NOT NULL,
    hospital_pk TEXT NOT NULL REFERENCES hospital(hospital_pk)
);
```

Design Decisions for "locations" table

The table for "locations" is designed to store geographic information associated with hospitals. We use zipcode as the primary key to represent each ZIP code once within our database and avoid repeating city, state, and county information across multiple hospital records. While multiple hospitals can share the same ZIP code in practice, this design choice keeps our schema simple and still allows us to associate hospitals with their regional attributes efficiently. Storing state, city, and fips_code with each ZIP supports geographic analysis without duplicating location attributes in the "hospital" table. ZIP codes are stored as text to preserve leading zeros and match real-world postal formatting, and the table structure remains easy to join from the "hospital" table through the foreign key reference.

```
In [ ]: %%sql
CREATE TABLE locations (
    zipcode INTEGER PRIMARY KEY,
    state TEXT NOT NULL,
    city TEXT NOT NULL,
    fips_code TEXT NOT NULL
);
```